



Housing Price Prediction by County

Jiayi Chen & Bell Luo for CSCI 0451

Motivation




Analyze the relationship between housing prices, demographic, and socioeconomic factors across counties



Inform and empower house buyers, sellers, renters, city planners, etc.

Goals

- **Combine** data from multiple sources to create a **comprehensive** dataset
- Create **interactive** maps to visualize housing price data by county
- Implement both **Classical Models** (including Linear Regression, SVM, Random Forest) and **Neural Networks** (through PyTorch) to analyze data



Data Sources

Zillow

- Typical home value
- By county & more
- Real-estate website
- Time series ready

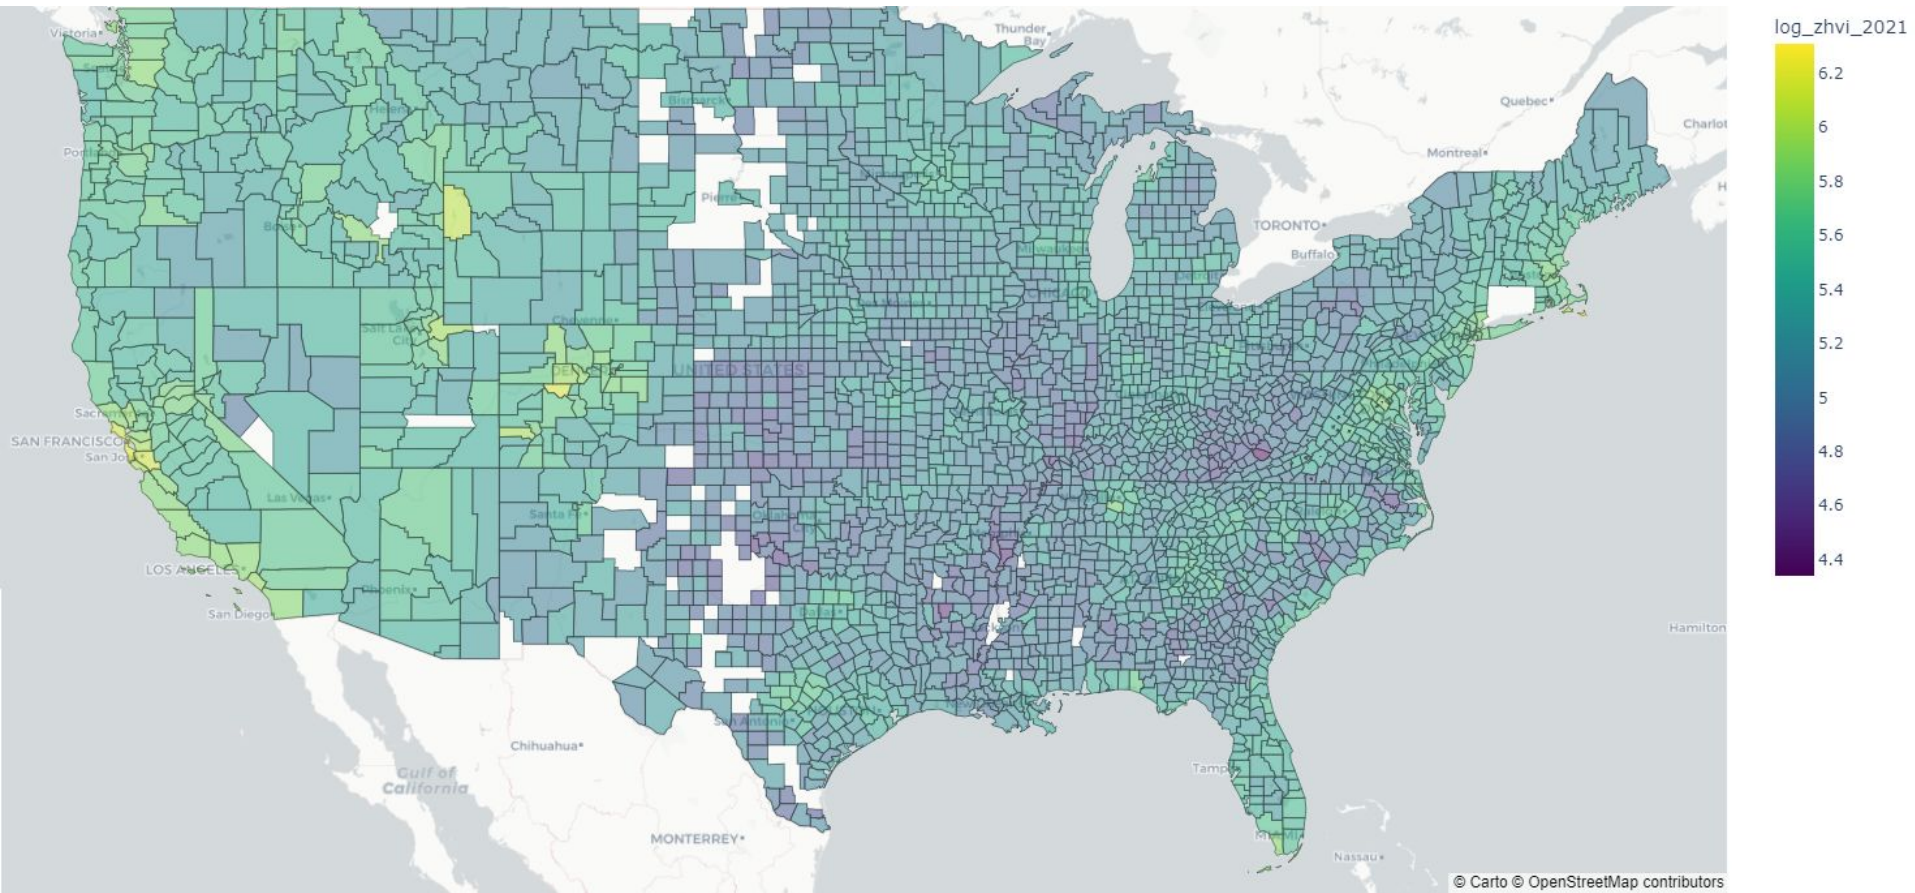
US Census


- ACS 5-Year Data
- Many predictors
- Population, income, education, etc.
- By county
- Time series ready

Miscellaneous

- More explanatory predictors
- Land area, Geospatial (longitude & latitude)
- By county

U.S. County-Level Housing Prices in 2021





Classical Models



Gradient Boosting

```
reg = GradientBoostingRegressor(**params)
reg.fit(X_train, np.ravel(y_train))
```

```
print(f"Score: {reg.score(X_test,
y_test):.4g}")
y_pred = reg.predict(X_test)
print(f"MSE: {mean_squared_error(y_test,
y_pred):.5e}")
```


```
Score: 0.863
MSE: 2.47022e+09
```


Random Forest

```
regr = make_pipeline(StandardScaler(),  
RandomForestRegressor(max_depth=20,  
random_state=42))  
regr.fit(X_train, np.ravel(y_train))  
print(f"Score: {regr.score(X_test,  
y_test):.4g}")  
y_pred = regr.predict(X_test)  
print(f"MSE: {mean_squared_error(y_test,  
y_pred):.5e}")
```

Score: 0.9109

MSE: 1.60781e+09






Neural Networks

Model Setup


```
class LinearRegression(nn.Module):  
    def __init__(self, input_size, hidden_size,  
                  output_size, num_hidden_layers):  
        super(LinearRegression, self).__init__()  
        self.input_layer = nn.Linear(input_size, hidden_size)  
        self.output_layer = nn.Linear(hidden_size,  
                                       output_size)  
  
        self.hidden_layers = nn.ModuleList()  
        for _ in range(num_hidden_layers):  
            self.hidden_layers.append(  
                nn.Linear(hidden_size, hidden_size))  
  
        self.activation = nn.ReLU()
```




Training Loop

```
tlh, vlh = training_loop(model,  
train_loader, test_loader, 1e-4, 500,  
l1_lambda=1e-4)
```

```
epoch 000: val_loss = 1.76887e+09  
epoch 050: val_loss = 2.88288e+08  
epoch 100: val_loss = 2.92384e+08  
epoch 150: val_loss = 2.78924e+08  
epoch 200: val_loss = 2.75191e+08  
epoch 250: val_loss = 2.73098e+08  
epoch 300: val_loss = 2.62659e+08  
epoch 350: val_loss = 2.76653e+08  
epoch 400: val_loss = 2.59071e+08
```






Further Tuning Needed

```
print(f"Score: {r2_score(y_test,  
y_pred):.4g}")  
print(f"MSE:  
{mean_squared_error(y_test,  
y_pred):.5e}")
```


Score: 0.5231
MSE: 8.60080e+09




Findings

Mean Decrease in Impurity (Importance)

MDI for Gradient Boosting



MDI for Random Forest



Important Features

Ordered by decreasing MDI:


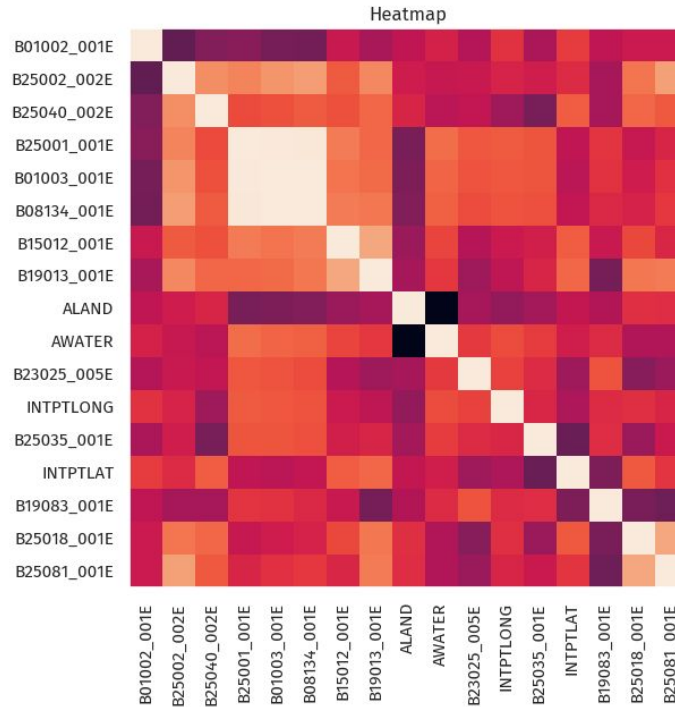

- **B19013_001E**: Median household income
- **INTPLONG**: Longitude
- **B15012_001E**: Count of Bachelor's degrees
- **B25035_001E**: Median year of construction

Model Performance

Ordered by decreasing test score:


- Random Forest Regression
- Gradient Boosting Regression
- Support Vector Regression
- Ridge Regression
- Linear Regression

Multicollinear or Correlated Features



Future Work

- Time series with RNN
 - Past prices
 - Changes in ACS variables
 - Fed rates & inflation
- Auditing bias
 - Across race & other variables
 - Cross-sectional
 - Temporal
- Canada & México?



Thank you

